

The scorecard's 61 tags cost nothing on the ranker and four gates on the gate

Results

The cascade refit on the GAIA scorecard's tag taxonomy — 60 tags from Scorecard - GAIA Catalog(Rasool-PII-PCI-PHI).csv, less routing_number, plus swift_code — scored on the eight sealed data/2-eval corpora under the unchanged evaluator.

	cascade_scorecard61	for reference
labels	61	58 before
thresholds selected	61 of 61	57 of 58 before — the one disabled head was routing_number, now dropped
micro F1 (<i>the ranker</i>)	0.7299	
micro precision	0.6360	
macro F0.5	0.6230	
macro F2	0.6651	target 0.6641, ceiling 0.6754
macro precision / recall	0.6244 / 0.6773	
worst measurable priority recall	0.6267	
prediction rate	0.8181	
document recall / precision / specificity	0.7975 / 0.8956 / 0.8792	
one-core p95	4.0290 ms	budget 5 ms — PASS

The target was met. Macro F2 0.6651 against a target of 0.6641 and a measured ceiling of 0.6754, so changing to the scorecard taxonomy cost the ranker nothing detectable. The document-level numbers are identical to the 58-label lineage to four decimals, which is expected rather than surprising: the gate is upstream of every tag head and its refit is unchanged.

The four restored tags, sealed

given_name, family_name, middle_name and street_number_and_name were folded into full_name and address in the previous lineage. Restored, all four cleared the recall floor on the training carve, and on the sealed corpora they look like this — precision first, since that is what ranks:

tag	corpus	support	P	R	F0.5
given_name	20000_pii_holdout	8,892	0.8367	0.9701	0.8603
	38937_openpii	29,335	0.7633	0.9974	0.8009

tag	corpus	support	P	R	F0.5
	30000_pii2_eval	5,453	0.6561	0.9118	0.6951
	10626_ai4privacy	3,501	0.3709	0.9797	0.4235
	10360_betterdataai	1,072	0.1382	0.7220	0.1649
family_name	20000_pii_holdout	8,337	0.7555	0.9745	0.7911
	38937_openpii	26,552	0.6901	0.9977	0.7355
	10360_betterdataai	247	0.0426	0.6964	0.0525
middle_name	20000_pii_holdout	6,039	0.9660	0.9631	0.9654
	38937_openpii	29,335	0.8047	0.9907	0.8361
	30000_pii2_eval	2,288	0.5099	0.9135	0.5593
	10360_betterdataai	14	0.0036	0.0714	0.0045
street_number_and_name	20000_pii_holdout	6,222	0.7757	0.9724	0.8084
	38937_openpii	15,267	0.6373	0.9954	0.6867
	10360_betterdataai	254	0.1251	0.7362	0.1500

Recall is 0.87–0.998 everywhere except betterdataai. Precision splits the corpora sharply, and the split is the one the feasibility probe predicted from gold alone before anything was fitted.

full_name **was not damaged by losing its subtypes**: 0.9739 / 0.9020 / 0.9331 precision on pii_holdout / openpii / betterdataai. Restoring the subtypes took nothing away from the parent tag.

Both declared policies

policy	verdict
headline — recall gates, ranks micro F1 → micro precision → macro F0.5 → macro F2	blocked , 25/55 priority gates, fails at password@10360_betterdataai
contra-view — document precision/specificity/recall + tag recall ≥ 0.75	blocked , doc constraints fail on govdocs2, 44/55 tag gates
one-core p95 ≤ 5 ms	PASS on both

NO FEASIBLE ARM. Nothing promoted, nothing packaged — the same 25/55 and the same failing scope as the 58-label v4, so the taxonomy change neither created nor fixed a gate failure.

TL;DR

- **61 labels**: the scorecard's 60, less sensitive_pci_routing_number (1 training row, 0 evaluation rows — the head that was already disabled), plus sensitive_pci_swift_code (542 / 130 rows, absent from the scorecard but a working detector).

- **Every scorecard tag has gold.** The coverage check found 0 tags with no gold anywhere, and exactly one sensitive tag in the gold that is not on the scorecard — `routing_number`, dropped deliberately.
- **The taxonomy change cost the ranker nothing:** macro F2 0.6651 against a 0.6641 target and a 0.6754 ceiling; document-level metrics unchanged to four decimals.
- **All 61 thresholds were selectable** — none fell below the support floor, against 57 of 58 before.
- **The restored tags work where the gold is consistent** and fail where it is not, exactly as predicted before fitting. `middle_name` reaches P 0.9660 on `pii_holdout` and P 0.0036 on `betterdataai`, whose name rows are 85.5% `full_name`-only.
- **Still blocked**, 25/55 priority gates, on `password@betterdataai` — the same scope that blocked the previous lineage. This run did not touch that.
- **Ranking now leads with precision.** On instruction, the headline policy ranks micro F1 → micro precision → macro F0.5 → macro F2. Recall was not demoted; it is the hard gate and always was.

Date	2026-08-26
Project	projects/pii-scorecard-60
Run	26-08-26-1
Request	"rebalance this model based on these 60 pii tags" → label space rebuilt to the scorecard; then "micro_f1 and precision to be the focus, but still have high recall, F@3"
Scope	1 model × 8 sealed corpora; 61-label catalogue; no hyperparameter search
CPU budget	1 core for the latency; training used the machine
Feasibility	plausible, ceiling 0.6754 vs target 0.6641 — <code>feasibility.json</code>
Outcome	target met on the new taxonomy; blocked by the recall gate; nothing promoted
Approval	<code>approvals/taxonomy-rebuild-61.json</code>

Model scoreboard

Model	Dataset	F1 @5	F1 @3	F1 @1	Macro F1	Docs/sec	ms/doc
pii-cascade-scorecard61	10360_betterdataai_ner_silver_eval_10.36k	0.225	0.180	0.126	0.263	248.2	4.0
pii-cascade-scorecard61	10626_ai4privacy_pii_masking_eval_10.63k	0.758	0.711	0.611	0.715	248.2	4.0
pii-cascade-scorecard61	20000_pii_holdout_20.00k	0.955	0.947	0.922	0.739	248.2	4.0
pii-cascade-scorecard61	30000_pii2_eval_25.15k	0.926	0.911	0.862	0.761	248.2	4.0

Model	Dataset	F1 @5	F1 @3	F1 @1	Macro F1	Docs/s ec	ms/ doc
pii-cascade-scorecard61	38937_openpii_pii_eval_38.94k	0.992	0.988	0.979	0.692	248.2	4.0
pii-cascade-scorecard61	4000_datax-dualjudge-evalset-1.32k	0.552	0.463	0.311		248.2	4.0
pii-cascade-scorecard61	5617_nemotron_eval_5.36k					248.2	4.0
pii-cascade-scorecard61	6589_govdocs2-dualjudge-eval20-3.53k	0.578	0.542	0.471		248.2	4.0

Measured 2026-08-26 · n=10,360, n=10,626, n=20,000, n=30,000, n=38,937, n=4,000, n=5,617, n=6,589.